



# MASCHINELLES LERNEN - UEBUNG 00

---

Sebastian Buschjäger

Technische Universität Dortmund - Fakultät Informatik - Lehrstuhl 8

October 22, 2018



## Structure of this course

### Goals

- Learn the basics of Applied Machine Learning
- Offer a place to discuss questions
- Prepare you for the exam



## Structure of this course

### Goals

- Learn the basics of Applied Machine Learning
- Offer a place to discuss questions
- Prepare you for the exam

### Structure

- ▶ There are no requirements to enter the exam (keine Studienleistung)
- ▶ There will be weekly exercises
- ▶ We will (probably) not publish sample solutions
- ▶ I will try to give a quick recap about the lecture each week
- ▶ We discuss what you deem necessary / interesting



## Recap: Notation

**Note** The input space can be (nearly) everything

**Make things easier** focus on  $d$ -dimensional vectors  $\vec{x} \in \mathcal{X} \subseteq \mathbb{R}^n$

$\mathcal{D}$	Feature 1	Feature 2	...	Feature d	Label
Example 1	$x_{11}$	$x_{12}$	...	$x_{1d}$	$y_1$
Example 2	$x_{21}$	$x_{22}$	...	$x_{2d}$	$y_2$
$\vdots$	$\vdots$	$\vdots$	$\ddots$	$\vdots$	$\vdots$
Example N	$x_{N1}$	$x_{N1}$	...	$x_{Nd}$	$y_N$



## Recap: Notation

**Note** The input space can be (nearly) everything

**Make things easier** focus on  $d$ -dimensional vectors  $\vec{x} \in \mathcal{X} \subseteq \mathbb{R}^n$

$\mathcal{D}$	Feature 1	Feature 2	...	Feature d	Label
Example 1	$x_{11}$	$x_{12}$	...	$x_{1d}$	$y_1$
Example 2	$x_{21}$	$x_{22}$	...	$x_{2d}$	$y_2$
$\vdots$	$\vdots$	$\vdots$	$\ddots$	$\vdots$	$\vdots$
Example N	$x_{N1}$	$x_{N1}$	...	$x_{Nd}$	$y_N$

Matrix  $X \in \mathbb{R}^{d \times N}$

Vector  $\vec{y} \in \mathcal{Y}^N$

**then** in short  $\mathcal{D} = (X, \vec{y})$



## Recap: Problems (1)

### Supervised Learning

- ▶ **Given** Set of labeled training examples / data  $\mathcal{D} = \{(\vec{x}_1, y_1), \dots, (\vec{x}_N, y_N) \mid (\vec{x}_i, y_i) \in \mathcal{X} \times \mathcal{Y}\}$
- ▶ **Find** A model  $f : \mathcal{X} \rightarrow \mathcal{Y}$  so that  $f(\vec{x}_i) \approx y_i$



## Recap: Problems (1)

### Supervised Learning

- ▶ **Given** Set of labeled training examples / data  $\mathcal{D} = \{(\vec{x}_1, y_1), \dots, (\vec{x}_N, y_N) \mid (\vec{x}_i, y_i) \in \mathcal{X} \times \mathcal{Y}\}$
- ▶ **Find** A model  $f : \mathcal{X} \rightarrow \mathcal{Y}$  so that  $f(\vec{x}_i) \approx y_i$

**Note 1** If  $|\mathcal{Y}| = 2$  its called binary classification

**Note 2** If  $\mathcal{Y} = \mathbb{R}$  its called regression

**Example** Classify songs into pop music vs. rap music



## Recap: Problems (2)

### Unsupervised Learning

- ▶ **Given** Set of un-labeled training data  $\mathcal{D} = \{\vec{x}_1, \dots, \vec{x}_N\}$
- ▶ **Find** A plausible model  $M$  that explains the data well

**Example** Put songs into similar groups to find new songs





## Recap: Problems (2)

### Unsupervised Learning

- ▶ **Given** Set of un-labeled training data  $\mathcal{D} = \{\vec{x}_1, \dots, \vec{x}_N\}$
- ▶ **Find** A plausible model  $M$  that explains the data well

**Example** Put songs into similar groups to find new songs

### Semi-Supervised Learning

- ▶ **Given** Set of partially-labeled training data  $\mathcal{D} = \{\vec{x}_1, \dots, \vec{x}_m, (\vec{x}_{m+1}, y_{m+1}), \dots, (\vec{x}_n, y_n)\}$
- ▶ **Find** A plausible model  $M$  that explains the data well with  $M(\vec{x}_i) \approx y_i$  for known  $y_i$

**Example** Put songs into similar groups knowing that some belong together / should never be in the same group



## Recap: Problems (3)

### Active Learning

- ▶ **Given** Set of partially-labeled training data  $\mathcal{D} = \{\vec{x}_1, \dots, \vec{x}_m, (\vec{x}_{m+1}, y_{m+1}), \dots, (\vec{x}_n, y_n)\}$
- ▶ **Find**  $\vec{x}_i \in \mathcal{D}$ , so that we acquire the true  $y_i$

**Example** For which songs should I hire an expert to classify them?



## Recap: Problems (3)

### Active Learning

- ▶ **Given** Set of partially-labeled training data  $\mathcal{D} = \{\vec{x}_1, \dots, \vec{x}_m, (\vec{x}_{m+1}, y_{m+1}), \dots, (\vec{x}_n, y_n)\}$
- ▶ **Find**  $\vec{x}_i \in \mathcal{D}$ , so that we acquire the true  $y_i$

**Example** For which songs should I hire an expert to classify them?

### Reinforcement Learning

- ▶ **Given** A set of actions  $A$  and a reward function  $r$
- ▶ **Goal** Find series of actions, so that reward is maximized

**Example** Maximize time the user listens to music by proposing new/similar songs



## Recap: Principles of Machine Learning

**Occam's Razor** Favor simple models before complex ones



## Recap: Principles of Machine Learning

**Occam's Razor** Favor simple models before complex ones

**No free lunch** All methods are equally good, but some favor certain data before other



## Recap: Principles of Machine Learning

**Occam's Razor** Favor simple models before complex ones

**No free lunch** All methods are equally good, but some favor certain data before other

**Bias-Variance trade-off** For squared error



## Recap: Principles of Machine Learning

**Occam's Razor** Favor simple models before complex ones

**No free lunch** All methods are equally good, but some favor certain data before other

**Bias-Variance trade-off** For squared error

Assume:  $t = f(x) + \varepsilon$  with  $\varepsilon \in \mathcal{N}(0, \sigma^2)$ , then

$$\begin{aligned}\mathbb{E}_{x,t,\mathcal{D}}[(t - f_{\mathcal{D}}(x))^2] &= (f(x) - \mathbb{E}_{\mathcal{D}|x}[f_{\mathcal{D}}(x)])^2 + \mathbb{V}_{\mathcal{D}|x}[f_{\mathcal{D}}(x)] + \mathbb{V}_{t|x}[\varepsilon] \\ &= (\text{bias})^2 + \text{variance} + \text{noise}\end{aligned}$$



## Extra: Measure Model quality

**Question** So, what is model quality?





## Extra: Measure Model quality

**Question** So, what is model quality?

1. how well explains the model training data?
2. can we give any guarantees for new predictions?
3. how well generalizes the model to new and unseen data?



## Extra: Measure Model quality (2)

**Fact** In binary classification we have two choices: predict 0 or 1  
→ 2 possible wrong predictions and 2 possible correct predictions



## Extra: Measure Model quality (2)

**Fact** In binary classification we have two choices: predict 0 or 1  
→ 2 possible wrong predictions and 2 possible correct predictions

**Visualization** Confusion matrix

	Predicted value	
True value	True positive (TP)	False negative (FN)
	False positive (FP)	True negative (TN)



## Extra: Measure Model quality (2)

**Fact** In binary classification we have two choices: predict 0 or 1  
→ 2 possible wrong predictions and 2 possible correct predictions

**Visualization** Confusion matrix

	Predicted value	
True value	True positive (TP)	False negative (FN)
	False positive (FP)	True negative (TN)

**Accuracy**  $Acc = \frac{TP+TN}{N}$

**Big Remark** The accuracy only tells us something about the data  $\mathcal{D}$  we know! There are no guarantees for new data



## Extra: Measure Model quality (3)

**Obviously** The best model has  $Acc = 1$ , the worst has  $Acc = 0$

**Observation** If we store all the data for look-up, then  $Acc = 1$



## Extra: Measure Model quality (3)

**Obviously** The best model has  $Acc = 1$ , the worst has  $Acc = 0$

**Observation** If we store all the data for look-up, then  $Acc = 1$

**Question** Is that what we want?

**Clear** This is just memorizing the training data, no real learning!

**Question** How well deals our model with new, yet unseen data?



## Extra: Measure Model quality (3)

**Obviously** The best model has  $\text{Acc} = 1$ , the worst has  $\text{Acc} = 0$

**Observation** If we store all the data for look-up, then  $\text{Acc} = 1$

**Question** Is that what we want?

**Clear** This is just memorizing the training data, no real learning!

**Question** How well deals our model with new, yet unseen data?

**Idea** Split data into training  $\mathcal{D}_{\text{Train}}$  and test data  $\mathcal{D}_{\text{Test}}$

**Then**  $\mathcal{D}_{\text{Test}}$  is new to the model  $f$

**Question** How to split  $\mathcal{D}$  ?



## Extra: Measure Model quality (4)

- 1) Test/Train Split** Split  $\mathcal{D}$  by size, e.g. 80% training and 20% test data
- Fast and easy to compute, but sensitive for “bad” splits.
  - Model quality might be over- or under-estimated





## Extra: Measure Model quality (4)

**1) Test/Train Split** Split  $\mathcal{D}$  by size, e.g. 80% training and 20% test data

- Fast and easy to compute, but sensitive for “bad” splits.
- Model quality might be over- or under-estimated

**2) Leave-One-Out** Use every example once for testing and train model on the remaining data.  
Average results.

- $N$  models are computed, but insensitive for “bad” splits.
- Usually impractical



## Extra: Measure Model quality (4)

**1) Test/Train Split** Split  $\mathcal{D}$  by size, e.g. 80% training and 20% test data

- Fast and easy to compute, but sensitive for “bad” splits.
- Model quality might be over- or under-estimated

**2) Leave-One-Out** Use every example once for testing and train model on the remaining data. Average results.

- $N$  models are computed, but insensitive for “bad” splits.
- Usually impractical

**3) K-fold cross validation** Split data into  $k$  buckets. Use every bucket once for testing / train model on the rest. Average results.

- Insensitive for “bad” splits and practical. Usually  $k \in \{5, 10\}$ .



## Extra: K nearest neighbor (K-NN) method

**Goal** Solve supervised problem

**Thus** We want a prediction method  $\hat{f}(\vec{x})$

**Observation** Examples  $\vec{x}_i$  and  $\vec{x}_j$  which are similar probably have the same label  $y_i = y_j$



## Extra: K nearest neighbor (K-NN) method

**Goal** Solve supervised problem

**Thus** We want a prediction method  $\hat{f}(\vec{x})$

**Observation** Examples  $\vec{x}_i$  and  $\vec{x}_j$  which are similar probably have the same label  $y_i = y_j$

**Idea** Given new and unseen observation  $\vec{x}$

- ▶ use distance function  $dist: \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$
- ▶ calculate  $d(\vec{x}, \vec{x}_i)$  for all  $i = 1, \dots, N$
- ▶ find  $k$  nearest neighbors of  $\vec{x}$   $S = \{(\vec{x}_1, y_1), \dots, (\vec{x}_k, y_k)\}$
- ▶ predict most common label in  $S$



## Extra: K nearest neighbor (K-NN) method

**Goal** Solve supervised problem

**Thus** We want a prediction method  $\hat{f}(\vec{x})$

**Observation** Examples  $\vec{x}_i$  and  $\vec{x}_j$  which are similar probably have the same label  $y_i = y_j$

**Idea** Given new and unseen observation  $\vec{x}$

- ▶ use distance function  $dist: \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$
- ▶ calculate  $d(\vec{x}, \vec{x}_i)$  for all  $i = 1, \dots, N$
- ▶ find  $k$  nearest neighbors of  $\vec{x}$   $S = \{(\vec{x}_1, y_1), \dots, (\vec{x}_k, y_k)\}$
- ▶ predict most common label in  $S$

**Note** If  $S$  has equal number of positive and negative examples, take a random class



## Extra: K-NN (Some Notes)

**Note** K-NN is a simple and large models. It simply stores  $\mathcal{D}$ .



## Extra: K-NN (Some Notes)

**Note** K-NN is a simple and large models. It simply stores  $\mathcal{D}$ .

**K-NN** has two parameters

- *dist* Models the distance of neighbors. This must fit the data given! Usually euclidean norm is a good start:

$$\text{dist}(\vec{x}_i, \vec{x}_j) = \sqrt{(\vec{x}_i - \vec{x}_j)^T \cdot (\vec{x}_i - \vec{x}_j)}$$

- *K* Models the number of neighbors we want to look at.



## Extra: K-NN (Some Notes)

**Note** K-NN is a simple and large models. It simply stores  $\mathcal{D}$ .

**K-NN** has two parameters

- *dist* Models the distance of neighbors. This must fit the data given! Usually euclidean norm is a good start:

$$\text{dist}(\vec{x}_i, \vec{x}_j) = \sqrt{(\vec{x}_i - \vec{x}_j)^T \cdot (\vec{x}_i - \vec{x}_j)}$$

- *K* Models the number of neighbors we want to look at.


**Note 2** K-NN can be used for regression as well. Just average the labels in  $S$  :

$$\hat{f}(\vec{x}) = \frac{1}{k} \sum_{y \in S} y$$






## Extra: K-NN Examples



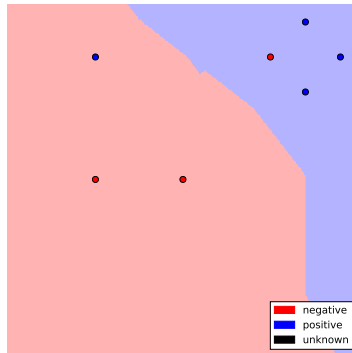
$k = 1$



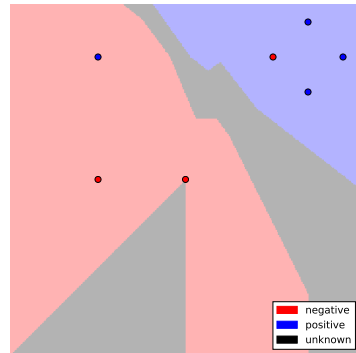
$k = 2$



## Extra: K-NN More examples



$k = 3$



$k = 4$



## Practical Data Science: Python

### Data Scientist use Python

- ▶ Dynamically scripting language
- ▶ Usually interpreted
- ▶ Only syntax is checked before execution (fails late)
- ▶ Indentation is a part of syntax (e.g. tabs instead of '{')



## Practical Data Science: Python

### Data Scientist use Python

- ▶ Dynamically scripting language
- ▶ Usually interpreted
- ▶ Only syntax is checked before execution (fails late)
- ▶ Indentation is a part of syntax (e.g. tabs instead of '{')

### Why Python?

- ▶ **Fast** Important algorithm have fast C/GPU backends
- ▶ **Flexible** Easy to test and try new things
- ▶ **Extensible** There are a lot of packages



## Practical Data Science: Packages

### Important Packages

- ▶ `numpy` Contains linear algebra functions
- ▶ `scipy` Contains common functions for scientific computing
- ▶ `matplotlib` Produces nice plots
- ▶ `pandas` Reads and handles files efficiently



## Practical Data Science: Packages

### Important Packages

- ▶ `numpy` Contains linear algebra functions
- ▶ `scipy` Contains common functions for scientific computing
- ▶ `matplotlib` Produces nice plots
- ▶ `pandas` Reads and handles files efficiently

### Important tools

- ▶ Python I recommend version  $\geq 3.5$
- ▶ `pip/anaconda` Install/Manage new packages
- ▶ `IPython` (Enhanced) Interactive python shell
- ▶ `Jupyter` Browser notebook for 'story' driven programming
- ▶ `plotly` Interactive plots on the web

**BUT** Your editor and raw python is your best tool